

neuMF_1m

April 1, 2025

```
[2]: import pandas as pd
import numpy as np

# Cargar el archivo ratings.dat especificando el separador '::'
ratings = pd.read_csv("../data/ml-1m/ml-1m/ratings.dat", sep="::",
    ↪engine="python", header=None,
    names=["UserID", "MovieID", "Rating", "Timestamp"])

# Mostrar las primeras filas para verificar la carga
print(ratings.head())
print(ratings.shape)
print(ratings.columns)

# Filtro: usuarios con al menos 5 ratings
user_counts = ratings['UserID'].value_counts()
ratings = ratings[ratings['UserID'].isin(user_counts[user_counts >= 5].index)]

# Filtro: películas con al menos 5 ratings
movie_counts = ratings['MovieID'].value_counts()
ratings = ratings[ratings['MovieID'].isin(movie_counts[movie_counts >= 5].
    ↪index)]

print(f"Usuarios después de filtrar: {ratings['UserID'].nunique()}")
print(f"Películas después de filtrar: {ratings['MovieID'].nunique()}")
```

```
   UserID  MovieID  Rating  Timestamp
0        1     1193        5  978300760
1        1      661        3  978302109
2        1      914        3  978301968
3        1     3408        4  978300275
4        1     2355        5  978824291
(1000209, 4)
Index(['UserID', 'MovieID', 'Rating', 'Timestamp'], dtype='object')
Usuarios después de filtrar: 6040
Películas después de filtrar: 3416
```

```
[3]: from sklearn.preprocessing import LabelEncoder
```

```

# Mapeo de IDs con LabelEncoder (más ordenado y reutilizable)
user_encoder = LabelEncoder()
movie_encoder = LabelEncoder()

ratings['userIndex'] = user_encoder.fit_transform(ratings['UserID'])
ratings['movieIndex'] = movie_encoder.fit_transform(ratings['MovieID'])

# Normalización a [0, 1]
ratings['rating_norm'] = ratings['Rating'] / 5.0

# Estandarización (media 0, desviación 1)
mean_rating = ratings['Rating'].mean()
std_rating = ratings['Rating'].std()
ratings['rating_std'] = (ratings['Rating'] - mean_rating) / std_rating

# Mostrar resumen
print(ratings[['Rating', 'rating_norm', 'rating_std']].head())
print(f"Media original: {mean_rating:.4f} | Desviación: {std_rating:.4f}")

```

	Rating	rating_norm	rating_std
0	5	1.0	1.269615
1	3	0.6	-0.521065
2	3	0.6	-0.521065
3	4	0.8	0.374275
4	5	1.0	1.269615

Media original: 3.5820 | Desviación: 1.1169

```

[4]: # Convertir timestamp a datetime
ratings['datetime'] = pd.to_datetime(ratings['Timestamp'], unit='s')

# Extraer año, mes y día de la semana
ratings['year'] = ratings['datetime'].dt.year
ratings['month'] = ratings['datetime'].dt.month
ratings['dayofweek'] = ratings['datetime'].dt.dayofweek # 0=Lunes, 6=Domingo

# Mostrar resumen
print(ratings[['Timestamp', 'datetime', 'year', 'month', 'dayofweek']].head())

```

	Timestamp	datetime	year	month	dayofweek
0	978300760	2000-12-31 22:12:40	2000	12	6
1	978302109	2000-12-31 22:35:09	2000	12	6
2	978301968	2000-12-31 22:32:48	2000	12	6
3	978300275	2000-12-31 22:04:35	2000	12	6
4	978824291	2001-01-06 23:38:11	2001	1	5

```

[5]: from sklearn.model_selection import train_test_split

# Elegir qué rating usar

```

```

rating_col = 'rating_norm' # cambia a 'rating_std' si quieres estandarizado

# Filtrar usuarios con al menos 3 ratings para hacer split por usuario
user_counts = ratings['UserID'].value_counts()
ratings_filtered = ratings[ratings['UserID'].isin(user_counts[user_counts >= 3].
    ↪index)]

# Split por usuario
train_list, val_list, test_list = [], [], []

for user_id, group in ratings_filtered.groupby('UserID'):
    user_train, user_temp = train_test_split(group, test_size=0.30,
    ↪random_state=42)
    user_val, user_test = train_test_split(user_temp, test_size=0.50,
    ↪random_state=42)

    train_list.append(user_train)
    val_list.append(user_val)
    test_list.append(user_test)

train_data = pd.concat(train_list).reset_index(drop=True)
val_data = pd.concat(val_list).reset_index(drop=True)
test_data = pd.concat(test_list).reset_index(drop=True)

print(f"Train size: {len(train_data)}")
print(f"Validation size: {len(val_data)}")
print(f"Test size: {len(test_data)}")

```

```

Train size: 697017
Validation size: 149779
Test size: 152815

```

```

[6]: import torch
from torch.utils.data import Dataset, DataLoader

# Convertir a tensores
train_user = torch.tensor(train_data['userIndex'].values, dtype=torch.long)
train_movie = torch.tensor(train_data['movieIndex'].values, dtype=torch.long)
train_rating = torch.tensor(train_data[rating_col].values, dtype=torch.float32)

val_user = torch.tensor(val_data['userIndex'].values, dtype=torch.long)
val_movie = torch.tensor(val_data['movieIndex'].values, dtype=torch.long)
val_rating = torch.tensor(val_data[rating_col].values, dtype=torch.float32)

test_user = torch.tensor(test_data['userIndex'].values, dtype=torch.long)
test_movie = torch.tensor(test_data['movieIndex'].values, dtype=torch.long)
test_rating = torch.tensor(test_data[rating_col].values, dtype=torch.float32)

```

```

# Dataset personalizado
class MovieLensDataset(Dataset):
    def __init__(self, users, movies, ratings):
        self.users = users
        self.movies = movies
        self.ratings = ratings

    def __len__(self):
        return len(self.ratings)

    def __getitem__(self, idx):
        return {
            'user': self.users[idx],
            'movie': self.movies[idx],
            'rating': self.ratings[idx]
        }

# Crear datasets
train_dataset = MovieLensDataset(train_user, train_movie, train_rating)
val_dataset = MovieLensDataset(val_user, val_movie, val_rating)
test_dataset = MovieLensDataset(test_user, test_movie, test_rating)

```

```

[7]: from torch.utils.data import DataLoader

batch_size = 512 # Puedes ajustar esto según tu GPU

train_loader = DataLoader(
    train_dataset,
    batch_size=batch_size,
    shuffle=True,
    num_workers=4,
    pin_memory=True
)

val_loader = DataLoader(
    val_dataset,
    batch_size=batch_size,
    shuffle=False,
    num_workers=4,
    pin_memory=True
)

test_loader = DataLoader(
    test_dataset,
    batch_size=batch_size,
    shuffle=False,

```

```

num_workers=4,
pin_memory=True
)

```

```

[8]: import torch.nn as nn
import torch.nn.functional as F

class NeuMF(nn.Module):
    def __init__(self, num_users, num_movies, embedding_dim_gmf=32,
↪embedding_dim_mlp=32, dropout=0.3):
        super().__init__()

        # Embeddings GMF
        self.user_embedding_gmf = nn.Embedding(num_users, embedding_dim_gmf)
        self.movie_embedding_gmf = nn.Embedding(num_movies, embedding_dim_gmf)

        # Embeddings MLP
        self.user_embedding_mlp = nn.Embedding(num_users, embedding_dim_mlp)
        self.movie_embedding_mlp = nn.Embedding(num_movies, embedding_dim_mlp)

        # MLP layers
        self.fc1 = nn.Linear(embedding_dim_mlp * 2, 128)
        self.bn1 = nn.BatchNorm1d(128)
        self.fc2 = nn.Linear(128, 64)
        self.bn2 = nn.BatchNorm1d(64)
        self.dropout = nn.Dropout(dropout)

        # Output layer: combina GMF + MLP
        self.output = nn.Linear(embedding_dim_gmf + 64, 1)

    def forward(self, user, movie):
        # GMF path
        user_gmf = self.user_embedding_gmf(user)
        movie_gmf = self.movie_embedding_gmf(movie)
        gmf_output = user_gmf * movie_gmf # Element-wise product

        # MLP path
        user_mlp = self.user_embedding_mlp(user)
        movie_mlp = self.movie_embedding_mlp(movie)
        mlp_input = torch.cat([user_mlp, movie_mlp], dim=1)
        x = F.relu(self.bn1(self.fc1(mlp_input)))
        x = self.dropout(x)
        x = F.relu(self.bn2(self.fc2(x)))
        x = self.dropout(x)

        # Combine GMF + MLP outputs
        final_input = torch.cat([gmf_output, x], dim=1)

```

```

        out = self.output(final_input)
        return out.squeeze()

```

```

[ ]: import torch.optim as optim
import torch.nn as nn

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f"Usando dispositivo: {device}")

# Aumentar dimensión de embeddings y añadir dropout adaptativo
model = NeuMF(
    num_users=len(user_encoder.classes_),
    num_movies=len(movie_encoder.classes_),
    embedding_dim_gmf=64, # ↑ de 32 a 64
    embedding_dim_mlp=64,
    dropout=0.5 # Regularización más agresiva
).to(device)

# Función de pérdida
criterion = nn.MSELoss()

# Optimizador
optimizer = optim.AdamW(model.parameters(), lr=0.002, weight_decay=1e-6)
scheduler = optim.lr_scheduler.OneCycleLR(
    optimizer,
    max_lr=0.005,
    steps_per_epoch=len(train_loader),
    epochs=30
)

```

Usando dispositivo: cuda

```

-----
ImportError                                Traceback (most recent call last)
Cell In[10], line 21
    18 criterion = nn.MSELoss()
    20 # Cambiar a optimizador LAMB con warmup
--> 21 from transformers.optimization import LAMB
    23 optimizer = LAMB(model.parameters(), lr=0.002, weight_decay=1e-6)
    24 scheduler = optim.lr_scheduler.OneCycleLR(
    25     optimizer,
    26     max_lr=0.005,
    27     steps_per_epoch=len(train_loader),
    28     epochs=30
    29 )

```

```
ImportError: cannot import name 'LAMB' from 'transformers.optimization' (/home/
↪eeguskiza/miniconda3/envs/work/lib/python3.10/site-packages/transformers/
↪optimization.py)
```

```
[ ]: num_epochs = 30
best_val_loss = float('inf')
patience = 5
early_stopping_counter = 0

for epoch in range(num_epochs):
    model.train()
    total_loss = 0

    for batch in train_loader:
        users = batch['user'].to(device)
        movies = batch['movie'].to(device)
        ratings = batch['rating'].to(device)

        preds = model(users, movies)
        loss = criterion(preds, ratings)

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        total_loss += loss.item() * len(ratings)

    avg_train_loss = total_loss / len(train_loader.dataset)

    # VALIDACIÓN
    model.eval()
    val_loss = 0
    with torch.no_grad():
        for batch in val_loader:
            users = batch['user'].to(device)
            movies = batch['movie'].to(device)
            ratings = batch['rating'].to(device)

            preds = model(users, movies)
            loss = criterion(preds, ratings)
            val_loss += loss.item() * len(ratings)

    avg_val_loss = val_loss / len(val_loader.dataset)
    scheduler.step(avg_val_loss)

    print(f"Epoch {epoch+1}/{num_epochs} | Train Loss: {avg_train_loss:.4f} |
↪Val Loss: {avg_val_loss:.4f}")
```

```

# EARLY STOPPING (opcional)
if avg_val_loss < best_val_loss:
    best_val_loss = avg_val_loss
    early_stopping_counter = 0
    torch.save(model.state_dict(), "best_model.pth")
else:
    early_stopping_counter += 1
    if early_stopping_counter >= patience:
        print(" Early stopping activado.")
        break

```

```

Epoch 1/30 | Train Loss: 0.0377 | Val Loss: 0.0308
Epoch 2/30 | Train Loss: 0.0310 | Val Loss: 0.0301
Epoch 3/30 | Train Loss: 0.0307 | Val Loss: 0.0299
Epoch 4/30 | Train Loss: 0.0306 | Val Loss: 0.0298
Epoch 5/30 | Train Loss: 0.0306 | Val Loss: 0.0299
Epoch 6/30 | Train Loss: 0.0306 | Val Loss: 0.0299
Epoch 7/30 | Train Loss: 0.0306 | Val Loss: 0.0300
Epoch 8/30 | Train Loss: 0.0296 | Val Loss: 0.0288
Epoch 9/30 | Train Loss: 0.0294 | Val Loss: 0.0286
Epoch 10/30 | Train Loss: 0.0293 | Val Loss: 0.0285
Epoch 11/30 | Train Loss: 0.0292 | Val Loss: 0.0284
Epoch 12/30 | Train Loss: 0.0291 | Val Loss: 0.0283
Epoch 13/30 | Train Loss: 0.0290 | Val Loss: 0.0282
Epoch 14/30 | Train Loss: 0.0290 | Val Loss: 0.0283
Epoch 15/30 | Train Loss: 0.0290 | Val Loss: 0.0282
Epoch 16/30 | Train Loss: 0.0290 | Val Loss: 0.0283
Epoch 17/30 | Train Loss: 0.0281 | Val Loss: 0.0273
Epoch 18/30 | Train Loss: 0.0279 | Val Loss: 0.0273
Epoch 19/30 | Train Loss: 0.0279 | Val Loss: 0.0273
Epoch 20/30 | Train Loss: 0.0278 | Val Loss: 0.0273
Epoch 21/30 | Train Loss: 0.0278 | Val Loss: 0.0273
Epoch 22/30 | Train Loss: 0.0278 | Val Loss: 0.0272
Epoch 23/30 | Train Loss: 0.0277 | Val Loss: 0.0272
Epoch 24/30 | Train Loss: 0.0277 | Val Loss: 0.0272
Epoch 25/30 | Train Loss: 0.0277 | Val Loss: 0.0272
Epoch 26/30 | Train Loss: 0.0277 | Val Loss: 0.0271
Epoch 27/30 | Train Loss: 0.0276 | Val Loss: 0.0272
Epoch 28/30 | Train Loss: 0.0276 | Val Loss: 0.0271
Epoch 29/30 | Train Loss: 0.0276 | Val Loss: 0.0271
Epoch 30/30 | Train Loss: 0.0276 | Val Loss: 0.0271

```

```

[ ]: import numpy as np
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

# Cargar el mejor modelo entrenado

```



```

model.eval()

y_true = []
y_pred = []

with torch.no_grad():
    for batch in test_loader:
        users = batch['user'].to(device)
        movies = batch['movie'].to(device)
        ratings = batch['rating'].to(device)

        preds = model(users, movies)

        # Si los ratings están normalizados, desnormalizamos
        preds = preds * 5
        ratings = ratings * 5

        y_true.extend(ratings.cpu().numpy())
        y_pred.extend(preds.cpu().numpy())

# Convertimos a arrays
y_true = np.array(y_true)
y_pred = np.array(y_pred)

# Cálculo de métricas
rmse = np.sqrt(mean_squared_error(y_true, y_pred))
mae = mean_absolute_error(y_true, y_pred)
r2 = r2_score(y_true, y_pred)

print(" MÉTRICAS DE ERROR - NeuMF")
print(f" RMSE: {rmse:.4f}")
print(f" MAE : {mae:.4f}")
print(f" R2 : {r2:.4f}")

```

```

MÉTRICAS DE ERROR - NeuMF
RMSE: 0.8239
MAE : 0.6384
R2 : 0.3970

```

```
[ ]: from collections import defaultdict
```

```

k = 10
user_preds = defaultdict(list)
user_truth = defaultdict(list)

model.eval()

```

```

with torch.no_grad():
    for batch in test_loader:
        users = batch['user'].to(device)
        movies = batch['movie'].to(device)
        ratings = batch['rating'].to(device)

        preds = model(users, movies)

        # Desnormalizamos (importante para comparar con escala original)
        preds = preds * 5
        ratings = ratings * 5

        for u, pred, true in zip(users.cpu().numpy(), preds.cpu().numpy(),
    ↪ ratings.cpu().numpy()):
            user_preds[u].append(pred)
            user_truth[u].append(true)

# Precision@K
precisions = []
for u in user_preds:
    preds_u = np.array(user_preds[u])
    truths_u = np.array(user_truth[u])

    top_k_indices = np.argsort(-preds_u)[:k]
    relevant = (truths_u >= 4.0)
    num_relevant = np.sum(relevant[top_k_indices])

    precision_u = num_relevant / k
    precisions.append(precision_u)

precision_at_k = np.mean(precisions)

# NDCG@K
def ndcg_at_k(relevances, k):
    relevances = np.asarray(relevances)[:k]
    if relevances.size == 0:
        return 0.0
    dcg = np.sum((2 ** relevances - 1) / np.log2(np.arange(2, relevances.size +
    ↪ 2)))
    ideal_relevances = np.sort(relevances)[::-1]
    idcg = np.sum((2 ** ideal_relevances - 1) / np.log2(np.arange(2,
    ↪ ideal_relevances.size + 2)))
    return dcg / idcg if idcg > 0 else 0.0

ndcgs = []
for u in user_preds:
    preds_u = np.array(user_preds[u])

```

```

truths_u = np.array(user_truth[u])
relevances = (truths_u >= 4.0).astype(int)
top_k_indices = np.argsort(-preds_u)[:k]
ndcg_u = ndcg_at_k(relevances[top_k_indices], k)
ndcgs.append(ndcg_u)

ndcg_at_k_value = np.mean(ndcgs)

print(f" Precision@{k}: {precision_at_k:.4f}")
print(f" NDCG@{k}: {ndcg_at_k_value:.4f}")

```

```

Precision@10: 0.5376
NDCG@10: 0.8899

```

```

[ ]: import time
import torch
import pandas as pd
import numpy as np
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from collections import defaultdict

# Cargar modelo entrenado
model.eval()

# Comenzar a contar el tiempo de evaluación
start_time = time.time()

# Evaluación sobre el set de test
y_true = []
y_pred = []

with torch.no_grad():
    for batch in test_loader:
        users = batch['user'].to(device)
        movies = batch['movie'].to(device)
        ratings = batch['rating'].to(device)

        preds = model(users, movies)

        # Desnormalizar si se entrenó en [0,1]
        preds = preds * 5
        ratings = ratings * 5

        y_true.extend(ratings.cpu().numpy())
        y_pred.extend(preds.cpu().numpy())

y_true = np.array(y_true)

```

```

y_pred = np.array(y_pred)

rmse = np.sqrt(mean_squared_error(y_true, y_pred))
mae = mean_absolute_error(y_true, y_pred)
r2 = r2_score(y_true, y_pred)

# Cálculo de Precision@10 y NDCG@10
k = 10
user_preds = defaultdict(list)
user_truth = defaultdict(list)

with torch.no_grad():
    for batch in test_loader:
        users = batch['user'].to(device)
        movies = batch['movie'].to(device)
        ratings = batch['rating'].to(device)
        preds = model(users, movies)

        preds = preds * 5
        ratings = ratings * 5

        for u, pred, true in zip(users.cpu().numpy(), preds.cpu().numpy(),
↪ ratings.cpu().numpy()):
            user_preds[u].append(pred)
            user_truth[u].append(true)

precisions = []
for u in user_preds:
    preds_u = np.array(user_preds[u])
    truths_u = np.array(user_truth[u])
    top_k_indices = np.argsort(-preds_u)[:k]
    relevant = (truths_u >= 4.0)
    num_relevant = np.sum(relevant[top_k_indices])
    precisions.append(num_relevant / k)
precision_at_k = np.mean(precisions)

def ndcg_at_k(relevances, k):
    relevances = np.asarray(relevances)[:k]
    if relevances.size == 0:
        return 0.0
    dcg = np.sum((2 ** relevances - 1) / np.log2(np.arange(2, relevances.size +
↪ 2)))
    ideal_relevances = np.sort(relevances)[::-1]
    idcg = np.sum((2 ** ideal_relevances - 1) / np.log2(np.arange(2,
↪ ideal_relevances.size + 2)))
    return dcg / idcg if idcg > 0 else 0.0

```

```

ndcgs = []
for u in user_preds:
    preds_u = np.array(user_preds[u])
    truths_u = np.array(user_truth[u])
    relevances = (truths_u >= 4.0).astype(int)
    top_k_indices = np.argsort(-preds_u)[:k]
    ndcgs.append(ndcg_at_k(relevances[top_k_indices], k))
ndcg_at_k_value = np.mean(ndcgs)

# Tiempo total de evaluación
eval_time = time.time() - start_time

# Resultados
print(" Evaluación final:")
print(f"RMSE           : {rmse:.4f}")
print(f"MAE           : {mae:.4f}")
print(f"R2           : {r2:.4f}")
print(f"Precision@{k}    : {precision_at_k:.4f}")
print(f"NDCG@{k}         : {ndcg_at_k_value:.4f}")
print(f"Eval Time (s)    : {eval_time:.2f}")

# Guardar en CSV
metrics = {
    "Model": "NeuMF (1M)",
    "Test RMSE": rmse,
    "Test MAE": mae,
    "Test R2": r2,
    "Precision@10": precision_at_k,
    "NDCG@10": ndcg_at_k_value,
    "Eval Time (s)": eval_time
}
metrics_df = pd.DataFrame([metrics])
#metrics_df.to_csv("neumf_10m_metrics_eval.csv", index=False)
print(" Métricas exportadas a 'neumf_1m_metrics_eval.csv'.")

```

```

Evaluación final:
RMSE           : 0.8239
MAE            : 0.6384
R2           : 0.3970
Precision@10   : 0.5376
NDCG@10        : 0.8899
Eval Time (s)  : 12.37
Métricas exportadas a 'neumf_1m_metrics_eval.csv'.

```