

ratings

April 1, 2025

0.1 Assignment 2

```
[1]: import torch
import numpy as np
torch.cuda.empty_cache()
```

Import necessary libraries and load the ratings

```
[3]: import pandas as pd

# Cargar el archivo ratings.csv
ratings = pd.read_csv('../data/ml-latest-small/ratings.csv')

# Ver las primeras filas
print(ratings.head())

# Revisar tamaño y columnas
print(ratings.shape)
print(ratings.columns)
```

```
   userId  movieId  rating  timestamp
0        1         1     4.0  964982703
1        1         3     4.0  964981247
2        1         6     4.0  964982224
3        1        47     5.0  964983815
4        1        50     5.0  964982931
(100836, 4)
Index(['userId', 'movieId', 'rating', 'timestamp'], dtype='object')
```

Preprocessing

```
[4]: # Filtrar usuarios con menos de 10 ratings
user_counts = ratings['userId'].value_counts()
ratings = ratings[ratings['userId'].isin(user_counts[user_counts >= 5].index)]

# Filtrar películas con menos de 10 ratings
movie_counts = ratings['movieId'].value_counts()
ratings = ratings[ratings['movieId'].isin(movie_counts[movie_counts >= 5].
↪index)]
```

```
print(f"Usuarios después de filtrar: {ratings['userId'].nunique()}")
print(f"Películas después de filtrar: {ratings['movieId'].nunique()}")
```

Usuarios después de filtrar: 610
Películas después de filtrar: 3650

```
[5]: # Obtener IDs únicos
unique_user_ids = ratings['userId'].unique()
unique_movie_ids = ratings['movieId'].unique()

print(f"Número de usuarios únicos: {len(unique_user_ids)}")
print(f"Número de películas únicas: {len(unique_movie_ids)}")

# Crear diccionarios de mapeo
userId_to_index = {user_id: idx for idx, user_id in enumerate(unique_user_ids)}
movieId_to_index = {movie_id: idx for idx, movie_id in
    ↪ enumerate(unique_movie_ids)}

# Aplicar el mapeo al DataFrame
ratings['userIndex'] = ratings['userId'].map(userId_to_index)
ratings['movieIndex'] = ratings['movieId'].map(movieId_to_index)

# Comprobar
print(ratings.head())
```

Número de usuarios únicos: 610
Número de películas únicas: 3650

	userId	movieId	rating	timestamp	userIndex	movieIndex
0	1	1	4.0	964982703	0	0
1	1	3	4.0	964981247	0	1
2	1	6	4.0	964982224	0	2
3	1	47	5.0	964983815	0	3
4	1	50	5.0	964982931	0	4

```
[6]: # Normalizamos ratings a [0, 1]
ratings['rating_norm'] = ratings['rating'] / 5.0
print(ratings[['rating', 'rating_norm']].head())
```

	rating	rating_norm
0	4.0	0.8
1	4.0	0.8
2	4.0	0.8
3	5.0	1.0
4	5.0	1.0

Split and Prepare

```
[7]: from sklearn.model_selection import train_test_split
```

```

# Primero filtramos usuarios con al menos 3 ratings
user_counts = ratings['userId'].value_counts()
ratings_filtered = ratings[ratings['userId'].isin(user_counts[user_counts >= 3].
    ↪index)]

# Luego aplicamos el split
train_list = []
val_list = []
test_list = []

for user_id, group in ratings_filtered.groupby('userId'):
    user_train, user_temp = train_test_split(group, test_size=0.30,
    ↪random_state=42)
    user_val, user_test = train_test_split(user_temp, test_size=0.50,
    ↪random_state=42)

    train_list.append(user_train)
    val_list.append(user_val)
    test_list.append(user_test)

train_data = pd.concat(train_list).reset_index(drop=True)
val_data = pd.concat(val_list).reset_index(drop=True)
test_data = pd.concat(test_list).reset_index(drop=True)

print(f"Train size: {len(train_data)}")
print(f"Validation size: {len(val_data)}")
print(f"Test size: {len(test_data)}")

```

Train size: 62932

Validation size: 13515

Test size: 13827

Dataloaders

```

[8]: import torch
from torch.utils.data import Dataset, DataLoader

# Convertir a tensores los índices de usuario, película y ratings
train_user = torch.tensor(train_data['userId'].values, dtype=torch.long)
train_movie = torch.tensor(train_data['movieIndex'].values, dtype=torch.long)
train_rating = torch.tensor(train_data['rating_norm'].values, dtype=torch.
    ↪float32)

val_user = torch.tensor(val_data['userId'].values, dtype=torch.long)
val_movie = torch.tensor(val_data['movieIndex'].values, dtype=torch.long)
val_rating = torch.tensor(val_data['rating_norm'].values, dtype=torch.float32)

test_user = torch.tensor(test_data['userId'].values, dtype=torch.long)

```

```
test_movie = torch.tensor(test_data['movieIndex'].values, dtype=torch.long)
test_rating = torch.tensor(test_data['rating_norm'].values, dtype=torch.float32)
```

Dataset Personalizado

```
[9]: class MovieLensDataset(Dataset):
    def __init__(self, users, movies, ratings):
        self.users = users
        self.movies = movies
        self.ratings = ratings

    def __len__(self):
        return len(self.ratings)

    def __getitem__(self, idx):
        return {
            'user': self.users[idx],
            'movie': self.movies[idx],
            'rating': self.ratings[idx]
        }
```

Dataloaders

```
[10]: batch_size = 512

train_dataset = MovieLensDataset(train_user, train_movie, train_rating)
val_dataset = MovieLensDataset(val_user, val_movie, val_rating)
test_dataset = MovieLensDataset(test_user, test_movie, test_rating)

train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=batch_size)
test_loader = DataLoader(test_dataset, batch_size=batch_size)
```

Simple Model

```
[11]: import torch.nn as nn
import torch.nn.functional as F

class NeuralCollaborativeFiltering(nn.Module):
    def __init__(self, num_users, num_movies, embedding_dim=64, dropout_rate=0.3):
        super(NeuralCollaborativeFiltering, self).__init__()

        self.user_embedding = nn.Embedding(num_users, embedding_dim)
        self.movie_embedding = nn.Embedding(num_movies, embedding_dim)

        # MLP con Dropout
        self.fc1 = nn.Linear(embedding_dim * 2, 128)
        self.fc2 = nn.Linear(128, 64)
```

```

        self.output_layer = nn.Linear(64, 1)

        self.dropout = nn.Dropout(dropout_rate)

    def forward(self, user, movie):
        user_embedded = self.user_embedding(user)
        movie_embedded = self.movie_embedding(movie)

        x = torch.cat([user_embedded, movie_embedded], dim=1)

        x = F.relu(self.fc1(x))
        x = self.dropout(x)
        x = F.relu(self.fc2(x))
        x = self.dropout(x)
        out = self.output_layer(x)

        return out.squeeze()

```

Optimizer

```

[12]: device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
      print(f"Usando dispositivo: {device}")

```

Usando dispositivo: cuda

```

[13]: num_users = len(userId_to_index)
      num_movies = len(movieId_to_index)

      model = NeuralCollaborativeFiltering(num_users=num_users, num_movies=num_movies).to(device)
      print("Modelo con Dropout creado correctamente")

```

Modelo con Dropout creado correctamente

```

[14]: import torch.optim as optim
      import torch.nn as nn

      criterion = nn.MSELoss()

      optimizer = optim.Adam(model.parameters(), lr=0.0005, weight_decay=1e-5)

```

Trainning

```

[34]: num_epochs = 10

      for epoch in range(num_epochs):
          model.train()
          total_loss = 0

          for batch in train_loader:

```

```

        users = batch['user'].to(device)
        movies = batch['movie'].to(device)
        ratings = batch['rating'].to(device)  # ← ya es rating_norm

        preds = model(users, movies)
        loss = criterion(preds, ratings)

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        total_loss += loss.item() * len(ratings)

    avg_train_loss = total_loss / len(train_loader.dataset)

    # Validation
    model.eval()
    val_loss = 0
    with torch.no_grad():
        for batch in val_loader:
            users = batch['user'].to(device)
            movies = batch['movie'].to(device)
            ratings = batch['rating'].to(device)

            preds = model(users, movies)
            loss = criterion(preds, ratings)
            val_loss += loss.item() * len(ratings)

    avg_val_loss = val_loss / len(val_loader.dataset)

    print(f"Epoch {epoch+1}: Train Loss = {avg_train_loss:.4f}, Val Loss = {avg_val_loss:.4f}")

```

```

Epoch 1: Train Loss = 0.0269, Val Loss = 0.0288
Epoch 2: Train Loss = 0.0265, Val Loss = 0.0287
Epoch 3: Train Loss = 0.0262, Val Loss = 0.0287
Epoch 4: Train Loss = 0.0259, Val Loss = 0.0287
Epoch 5: Train Loss = 0.0257, Val Loss = 0.0287
Epoch 6: Train Loss = 0.0252, Val Loss = 0.0288
Epoch 7: Train Loss = 0.0250, Val Loss = 0.0290
Epoch 8: Train Loss = 0.0246, Val Loss = 0.0291
Epoch 9: Train Loss = 0.0242, Val Loss = 0.0294
Epoch 10: Train Loss = 0.0238, Val Loss = 0.0294

```

```

[36]: torch.save(model.state_dict(), "solo_ratings.pth")
      print("Modelo guardado correctamente.")

```

Modelo guardado correctamente.

Model evaluation

RMSE

```
[26]: all_preds = []
      all_truth = []

      model.eval()
      with torch.no_grad():
          for batch in test_loader:
              users = batch['user'].to(device, dtype=torch.long)
              movies = batch['movie'].to(device, dtype=torch.long)
              ratings = batch['rating'].to(device, dtype=torch.float) # Asumo ↵
              ↪ ratings como float

              preds = model(users, movies)

              # Desnormalización (si ratings/preds están en 0-1)
              all_preds.extend((preds * 5).cpu().numpy())
              all_truth.extend((ratings * 5).cpu().numpy())

      all_preds = np.array(all_preds)
      all_truth = np.array(all_truth)

      rmse = np.sqrt(np.mean((all_preds - all_truth) ** 2))
```

MAE

```
[27]: mae = np.mean(np.abs(all_preds - all_truth))
```

R-Square

```
[28]: from sklearn.metrics import r2_score
      r2 = r2_score(all_truth, all_preds)
```

Precision

```
[29]: from collections import defaultdict

      k = 10
      user_preds = defaultdict(list)
      user_truth = defaultdict(list)

      model.eval()
      with torch.no_grad():
          for batch in test_loader:
              users = batch['user'].to(device, dtype=torch.long)
              movies = batch['movie'].to(device, dtype=torch.long)
              ratings = batch['rating'].to(device, dtype=torch.float)

              preds = model(users, movies)
```

```

    # Desnormalizamos a escala 0.5 - 5
    preds = preds * 5
    ratings = ratings * 5

    for u, pred, true in zip(users.cpu().numpy(), preds.cpu().numpy(),
↪ ratings.cpu().numpy()):
        user_preds[u].append(pred)
        user_truth[u].append(true)

# Calculamos Precision@K
precisions = []
for u in user_preds:
    preds_u = np.array(user_preds[u])
    truths_u = np.array(user_truth[u])

    # Ordenar por predicción descendente y coger top K
    top_k_indices = np.argsort(-preds_u)[:k]

    # Definir relevantes: ratings reales >= 4.0
    relevant = (truths_u >= 4.0)
    num_relevant = np.sum(relevant[top_k_indices])

    precision_u = num_relevant / k
    precisions.append(precision_u)

precision_at_k = np.mean(precisions)

```

NDCG@K

```

[30]: def ndcg_at_k(relevances, k):
    relevances = np.asarray(relevances)[:k]
    if relevances.size == 0:
        return 0.0
    # DCG:  $(2^{\text{rel}} - 1) / \log_2(\text{pos} + 1)$ 
    dcg = np.sum((2 ** relevances - 1) / np.log2(np.arange(2, relevances.size +
↪ 2)))

    # Ideal DCG: orden perfecto
    ideal_relevances = np.sort(relevances)[::-1]
    idcg = np.sum((2 ** ideal_relevances - 1) / np.log2(np.arange(2,
↪ ideal_relevances.size + 2)))

    return dcg / idcg if idcg > 0 else 0.0

# Calculamos NDCG@K para cada usuario
ndcgs = []
for u in user_preds:

```



```

preds_u = np.array(user_preds[u])
truths_u = np.array(user_truth[u])

# Relevancia binaria: 1 si rating real >= 4.0, 0 en caso contrario
relevances = (truths_u >= 4.0).astype(int)

# Ordenar por predicción descendente y coger top-K
top_k_indices = np.argsort(-preds_u)[:k]

ndcg_u = ndcg_at_k(relevances[top_k_indices], k)
ndcgs.append(ndcg_u)

# Media total
ndcg_at_k_value = np.mean(ndcgs)

```

All Metrics

```

[31]: from tabulate import tabulate

# Suponiendo que ya tienes las variables calculadas:
# rmse, mae, r2, precision_at_k, ndcg_at_k_value

# Aseguramos que todos los valores sean Python floats y estén redondeados
metrics = {
    'RMSE': float(np.round(rmse, 4)),
    'MAE': float(np.round(mae, 4)),
    'R2': float(np.round(r2, 4)),
    'Pre': float(np.round(precision_at_k, 4)),
    'NDCG@10': float(np.round(ndcg_at_k_value, 4))
}

# Crear un DataFrame
metrics_df = pd.DataFrame(list(metrics.items()), columns=['Métrica', 'Valor'])

# Alternativa: formatear explícitamente los valores en el DataFrame
metrics_df['Valor'] = metrics_df['Valor'].apply(lambda x: f"{x:.4f}")

import matplotlib.pyplot as plt

# Creamos figura y eje
fig, ax = plt.subplots(figsize=(6, 2))

# Ocultamos los ejes
ax.axis('off')

# Creamos la tabla visualmente
table = ax.table(cellText=metrics_df.values,
                 colLabels=metrics_df.columns,

```

```

        loc='center',
        cellLoc='center')

# Ajustes estéticos
table.auto_set_font_size(False)
table.set_fontsize(12)
table.scale(1.2, 1.2)

plt.title("Métricas del Modelo", fontsize=14, pad=20)
plt.show()

```

Métricas del Modelo

Métrica	Valor
RMSE	0.8688
MAE	0.6677
R2	0.3092
Pre	0.5180
NDCG@10	0.8535

```

[37]: import numpy as np
import pandas as pd
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from collections import defaultdict
import torch

# Cargar el modelo entrenado (ajusta la ruta según corresponda)
model.load_state_dict(torch.load("solo_rantings.pth", map_location=device))
model.eval()

# Evaluación en el conjunto de test
y_true = []
y_pred = []
with torch.no_grad():
    for batch in test_loader:
        users = batch['user'].to(device)
        movies = batch['movie'].to(device)
        ratings = batch['rating'].to(device)
        preds = model(users, movies)
        # Desnormalizamos a [0, 5]
        preds = preds * 5

```

```

        ratings = ratings * 5
        y_true.extend(ratings.cpu().numpy())
        y_pred.extend(preds.cpu().numpy())
y_true = np.array(y_true)
y_pred = np.array(y_pred)

# Métricas de error
rmse = np.sqrt(mean_squared_error(y_true, y_pred))
mae = mean_absolute_error(y_true, y_pred)
r2 = r2_score(y_true, y_pred)

# Cálculo de Precision@10 y NDCG@10
k = 10
user_preds = defaultdict(list)
user_truth = defaultdict(list)
with torch.no_grad():
    for batch in test_loader:
        users = batch['user'].to(device)
        movies = batch['movie'].to(device)
        ratings = batch['rating'].to(device)
        preds = model(users, movies)
        preds = preds * 5
        ratings = ratings * 5
        for u, pred, true in zip(users.cpu().numpy(), preds.cpu().numpy(),
↪ ratings.cpu().numpy()):
            user_preds[u].append(pred)
            user_truth[u].append(true)

precisions = []
for u in user_preds:
    preds_u = np.array(user_preds[u])
    truths_u = np.array(user_truth[u])
    top_k_indices = np.argsort(-preds_u)[:k]
    relevant = (truths_u >= 4.0)
    num_relevant = np.sum(relevant[top_k_indices])
    precisions.append(num_relevant / k)
precision_at_k = np.mean(precisions)

def ndcg_at_k(relevances, k):
    relevances = np.asarray(relevances)[:k]
    if relevances.size == 0:
        return 0.0
    dcg = np.sum((2**relevances - 1) / np.log2(np.arange(2, relevances.size +
↪ 2)))
    ideal_relevances = np.sort(relevances[::-1])
    idcg = np.sum((2**ideal_relevances - 1) / np.log2(np.arange(2,
↪ ideal_relevances.size + 2)))

```

```

        return dcg / idcg if idcg > 0 else 0.0

ndcgs = []
for u in user_preds:
    preds_u = np.array(user_preds[u])
    truths_u = np.array(user_truth[u])
    relevances = (truths_u >= 4.0).astype(int)
    top_k_indices = np.argsort(-preds_u)[:k]
    ndcgs.append(ndcg_at_k(relevances[top_k_indices], k))
ndcg_at_k_value = np.mean(ndcgs)

# Mostrar resultados en consola
print("Métricas de Test:")
print(f"RMSE: {rmse:.4f}")
print(f"MAE : {mae:.4f}")
print(f"R²  : {r2:.4f}")
print(f"Precision@{k}: {precision_at_k:.4f}")
print(f"NDCG@{k}: {ndcg_at_k_value:.4f}")

# Exportar las métricas a un CSV
metrics = {
    "Model": "NeuralCollaborativeFiltering (100K)",
    "Test RMSE": rmse,
    "Test MAE": mae,
    "Test R2": r2,
    "Precision@10": precision_at_k,
    "NDCG@10": ndcg_at_k_value
}
metrics_df = pd.DataFrame([metrics])
metrics_df.to_csv("ncf_simple_100k_metrics.csv", index=False)
print("Métricas exportadas a 'ncf_simple_100k_metrics.csv'.")

```

Métricas de Test:

RMSE: 0.8758

MAE : 0.6692

R² : 0.2980

Precision@10: 0.5125

NDCG@10: 0.8486

Métricas exportadas a 'ncf_simple_100k_metrics.csv'.