

ratings

April 3, 2025

0.1 Assignment 2

```
[16]: import torch
import numpy as np
torch.cuda.empty_cache()
```

Import necessary libraries and load the ratings

```
[1]: import pandas as pd

# Cargar el archivo ratings.csv
ratings = pd.read_csv('../data/ml-latest-small/ratings.csv')

# Ver las primeras filas
print(ratings.head())

# Revisar tamaño y columnas
print(ratings.shape)
print(ratings.columns)
```

```
   userId  movieId  rating  timestamp
0        1         1     4.0  964982703
1        1         3     4.0  964981247
2        1         6     4.0  964982224
3        1        47     5.0  964983815
4        1        50     5.0  964982931
(100836, 4)
Index(['userId', 'movieId', 'rating', 'timestamp'], dtype='object')
```

Preprocessing

```
[2]: # Filtrar usuarios con menos de 10 ratings
user_counts = ratings['userId'].value_counts()
ratings = ratings[ratings['userId'].isin(user_counts[user_counts >= 5].index)]

# Filtrar películas con menos de 10 ratings
movie_counts = ratings['movieId'].value_counts()
ratings = ratings[ratings['movieId'].isin(movie_counts[movie_counts >= 5].
↪index)]
```

```
print(f"Usuarios después de filtrar: {ratings['userId'].nunique()}")
print(f"Películas después de filtrar: {ratings['movieId'].nunique()}")
```

Usuarios después de filtrar: 610
Películas después de filtrar: 3650

```
[3]: # Obtener IDs únicos
unique_user_ids = ratings['userId'].unique()
unique_movie_ids = ratings['movieId'].unique()

print(f"Número de usuarios únicos: {len(unique_user_ids)}")
print(f"Número de películas únicas: {len(unique_movie_ids)}")

# Crear diccionarios de mapeo
userId_to_index = {user_id: idx for idx, user_id in enumerate(unique_user_ids)}
movieId_to_index = {movie_id: idx for idx, movie_id in
    ↪ enumerate(unique_movie_ids)}

# Aplicar el mapeo al DataFrame
ratings['userIndex'] = ratings['userId'].map(userId_to_index)
ratings['movieIndex'] = ratings['movieId'].map(movieId_to_index)

# Comprobar
print(ratings.head())
```

Número de usuarios únicos: 610
Número de películas únicas: 3650

	userId	movieId	rating	timestamp	userIndex	movieIndex
0	1	1	4.0	964982703	0	0
1	1	3	4.0	964981247	0	1
2	1	6	4.0	964982224	0	2
3	1	47	5.0	964983815	0	3
4	1	50	5.0	964982931	0	4

```
[4]: # Normalizamos ratings a [0, 1]
ratings['rating_norm'] = ratings['rating'] / 5.0
print(ratings[['rating', 'rating_norm']].head())
```

	rating	rating_norm
0	4.0	0.8
1	4.0	0.8
2	4.0	0.8
3	5.0	1.0
4	5.0	1.0

Split and Prepare

```
[5]: from sklearn.model_selection import train_test_split
```

```

# Primero filtramos usuarios con al menos 3 ratings
user_counts = ratings['userId'].value_counts()
ratings_filtered = ratings[ratings['userId'].isin(user_counts[user_counts >= 3].
    ↪index)]

# Luego aplicamos el split
train_list = []
val_list = []
test_list = []

for user_id, group in ratings_filtered.groupby('userId'):
    user_train, user_temp = train_test_split(group, test_size=0.30,
    ↪random_state=42)
    user_val, user_test = train_test_split(user_temp, test_size=0.50,
    ↪random_state=42)

    train_list.append(user_train)
    val_list.append(user_val)
    test_list.append(user_test)

train_data = pd.concat(train_list).reset_index(drop=True)
val_data = pd.concat(val_list).reset_index(drop=True)
test_data = pd.concat(test_list).reset_index(drop=True)

print(f"Train size: {len(train_data)}")
print(f"Validation size: {len(val_data)}")
print(f"Test size: {len(test_data)}")

```

Train size: 62932

Validation size: 13515

Test size: 13827

Dataloaders

```

[6]: import torch
from torch.utils.data import Dataset, DataLoader

# Convertir a tensores los índices de usuario, película y ratings
train_user = torch.tensor(train_data['userId'].values, dtype=torch.long)
train_movie = torch.tensor(train_data['movieIndex'].values, dtype=torch.long)
train_rating = torch.tensor(train_data['rating_norm'].values, dtype=torch.
    ↪float32)

val_user = torch.tensor(val_data['userId'].values, dtype=torch.long)
val_movie = torch.tensor(val_data['movieIndex'].values, dtype=torch.long)
val_rating = torch.tensor(val_data['rating_norm'].values, dtype=torch.float32)

test_user = torch.tensor(test_data['userId'].values, dtype=torch.long)

```

```
test_movie = torch.tensor(test_data['movieIndex'].values, dtype=torch.long)
test_rating = torch.tensor(test_data['rating_norm'].values, dtype=torch.float32)
```

Dataset Personalizado

```
[7]: class MovieLensDataset(Dataset):
    def __init__(self, users, movies, ratings):
        self.users = users
        self.movies = movies
        self.ratings = ratings

    def __len__(self):
        return len(self.ratings)

    def __getitem__(self, idx):
        return {
            'user': self.users[idx],
            'movie': self.movies[idx],
            'rating': self.ratings[idx]
        }
```

Dataloaders

```
[8]: batch_size = 512

train_dataset = MovieLensDataset(train_user, train_movie, train_rating)
val_dataset = MovieLensDataset(val_user, val_movie, val_rating)
test_dataset = MovieLensDataset(test_user, test_movie, test_rating)

train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=batch_size)
test_loader = DataLoader(test_dataset, batch_size=batch_size)
```

Simple Model

```
[9]: import torch.nn as nn
import torch.nn.functional as F

class NeuralCollaborativeFiltering(nn.Module):
    def __init__(self, num_users, num_movies, embedding_dim=64, dropout_rate=0.3):
        super(NeuralCollaborativeFiltering, self).__init__()

        self.user_embedding = nn.Embedding(num_users, embedding_dim)
        self.movie_embedding = nn.Embedding(num_movies, embedding_dim)

        # MLP con Dropout
        self.fc1 = nn.Linear(embedding_dim * 2, 128)
        self.fc2 = nn.Linear(128, 64)
```

```

        self.output_layer = nn.Linear(64, 1)

        self.dropout = nn.Dropout(dropout_rate)

    def forward(self, user, movie):
        user_embedded = self.user_embedding(user)
        movie_embedded = self.movie_embedding(movie)

        x = torch.cat([user_embedded, movie_embedded], dim=1)

        x = F.relu(self.fc1(x))
        x = self.dropout(x)
        x = F.relu(self.fc2(x))
        x = self.dropout(x)
        out = self.output_layer(x)

        return out.squeeze()

```

Optimizer

```

[10]: device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
       print(f"Usando dispositivo: {device}")

```

Usando dispositivo: cuda

```

[11]: num_users = len(userId_to_index)
       num_movies = len(movieId_to_index)

       model = NeuralCollaborativeFiltering(num_users=num_users, num_movies=num_movies).to(device)
       print("Modelo con Dropout creado correctamente")

```

Modelo con Dropout creado correctamente

```

[12]: import torch.optim as optim
       import torch.nn as nn

       criterion = nn.MSELoss()

       optimizer = optim.Adam(model.parameters(), lr=0.0005, weight_decay=1e-5)

```

Trainning

```

[13]: num_epochs = 10

       for epoch in range(num_epochs):
           model.train()
           total_loss = 0

           for batch in train_loader:

```

```

        users = batch['user'].to(device)
        movies = batch['movie'].to(device)
        ratings = batch['rating'].to(device)  # ← ya es rating_norm

        preds = model(users, movies)
        loss = criterion(preds, ratings)

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        total_loss += loss.item() * len(ratings)

    avg_train_loss = total_loss / len(train_loader.dataset)

    # Validation
    model.eval()
    val_loss = 0
    with torch.no_grad():
        for batch in val_loader:
            users = batch['user'].to(device)
            movies = batch['movie'].to(device)
            ratings = batch['rating'].to(device)

            preds = model(users, movies)
            loss = criterion(preds, ratings)
            val_loss += loss.item() * len(ratings)

    avg_val_loss = val_loss / len(val_loader.dataset)

    print(f"Epoch {epoch+1}: Train Loss = {avg_train_loss:.4f}, Val Loss = {avg_val_loss:.4f}")

```

```

Epoch 1: Train Loss = 0.0792, Val Loss = 0.0420
Epoch 2: Train Loss = 0.0526, Val Loss = 0.0398
Epoch 3: Train Loss = 0.0486, Val Loss = 0.0382
Epoch 4: Train Loss = 0.0458, Val Loss = 0.0371
Epoch 5: Train Loss = 0.0435, Val Loss = 0.0360
Epoch 6: Train Loss = 0.0422, Val Loss = 0.0355
Epoch 7: Train Loss = 0.0411, Val Loss = 0.0347
Epoch 8: Train Loss = 0.0400, Val Loss = 0.0347
Epoch 9: Train Loss = 0.0393, Val Loss = 0.0341
Epoch 10: Train Loss = 0.0386, Val Loss = 0.0338

```

```

[14]: torch.save(model.state_dict(), "solo_ratings.pth")
      print("Modelo guardado correctamente.")

```

Modelo guardado correctamente.

Model evaluation

RMSE

```
[18]: all_preds = []
      all_truth = []

      model.eval()
      with torch.no_grad():
          for batch in test_loader:
              users = batch['user'].to(device, dtype=torch.long)
              movies = batch['movie'].to(device, dtype=torch.long)
              ratings = batch['rating'].to(device, dtype=torch.float) # Asumo ↵
              ↵ ratings como float

              preds = model(users, movies)

              # Desnormalización (si ratings/preds están en 0-1)
              all_preds.extend((preds * 5).cpu().numpy())
              all_truth.extend((ratings * 5).cpu().numpy())

      all_preds = np.array(all_preds)
      all_truth = np.array(all_truth)

      rmse = np.sqrt(np.mean((all_preds - all_truth) ** 2))
```

MAE

```
[19]: mae = np.mean(np.abs(all_preds - all_truth))
```

R-Square

```
[20]: from sklearn.metrics import r2_score
      r2 = r2_score(all_truth, all_preds)
```

Precision

```
[21]: from collections import defaultdict

      k = 10
      user_preds = defaultdict(list)
      user_truth = defaultdict(list)

      model.eval()
      with torch.no_grad():
          for batch in test_loader:
              users = batch['user'].to(device, dtype=torch.long)
              movies = batch['movie'].to(device, dtype=torch.long)
              ratings = batch['rating'].to(device, dtype=torch.float)

              preds = model(users, movies)
```

```

    # Desnormalizamos a escala 0.5 - 5
    preds = preds * 5
    ratings = ratings * 5

    for u, pred, true in zip(users.cpu().numpy(), preds.cpu().numpy(),
    ↪ ratings.cpu().numpy()):
        user_preds[u].append(pred)
        user_truth[u].append(true)

# Calculamos Precision@K
precisions = []
for u in user_preds:
    preds_u = np.array(user_preds[u])
    truths_u = np.array(user_truth[u])

    # Ordenar por predicción descendente y coger top K
    top_k_indices = np.argsort(-preds_u)[:k]

    # Definir relevantes: ratings reales >= 4.0
    relevant = (truths_u >= 4.0)
    num_relevant = np.sum(relevant[top_k_indices])

    precision_u = num_relevant / k
    precisions.append(precision_u)

precision_at_k = np.mean(precisions)

```

NDCG@K

```

[22]: def ndcg_at_k(relevances, k):
    relevances = np.asarray(relevances)[:k]
    if relevances.size == 0:
        return 0.0
    # DCG:  $(2^{\text{rel}} - 1) / \log_2(\text{pos} + 1)$ 
    dcg = np.sum((2 ** relevances - 1) / np.log2(np.arange(2, relevances.size +
    ↪ 2)))

    # Ideal DCG: orden perfecto
    ideal_relevances = np.sort(relevances)[::-1]
    idcg = np.sum((2 ** ideal_relevances - 1) / np.log2(np.arange(2,
    ↪ ideal_relevances.size + 2)))

    return dcg / idcg if idcg > 0 else 0.0

# Calculamos NDCG@K para cada usuario
ndcgs = []
for u in user_preds:

```



```

preds_u = np.array(user_preds[u])
truths_u = np.array(user_truth[u])

# Relevancia binaria: 1 si rating real >= 4.0, 0 en caso contrario
relevances = (truths_u >= 4.0).astype(int)

# Ordenar por predicción descendente y coger top-K
top_k_indices = np.argsort(-preds_u)[:k]

ndcg_u = ndcg_at_k(relevances[top_k_indices], k)
ndcgs.append(ndcg_u)

# Media total
ndcg_at_k_value = np.mean(ndcgs)

```

Recall@10 y F1@10

[23]: # Bloque: Cálculo de Recall@10 y F1@10

```

k = 10
recalls = []
for u in user_preds:
    preds_u = np.array(user_preds[u])
    truths_u = np.array(user_truth[u])
    # Obtener los índices de los top k elementos con mayor predicción
    top_k_indices = np.argsort(-preds_u)[:k]
    # Definir items relevantes (ej.: rating real >= 4.0)
    relevant = (truths_u >= 4.0)
    total_relevant = np.sum(relevant)
    if total_relevant > 0:
        recall_u = np.sum(relevant[top_k_indices]) / total_relevant
        recalls.append(recall_u)
    else:
        recalls.append(0.0)
recall_at_k = np.mean(recalls)
# Calcular F1@10: media armónica de Precision@10 y Recall@10
f1_at_k = 2 * (precision_at_k * recall_at_k) / (precision_at_k + recall_at_k)
↪ if (precision_at_k + recall_at_k) > 0 else 0.0

```

[]: # Bloque: Consolidación de métricas en un CSV

```

# Se asume que ya tienes las siguientes variables calculadas:
# rmse, mae, r2, precision_at_k, ndcg_at_k_value, recall_at_k, f1_at_k

# Crear un diccionario con las métricas del modelo actual
current_metrics = {
    'Modelo': 'R (100K)', # Cambia este identificador si es necesario
    'Test RMSE': float(np.round(rmse, 4)),

```

```

    'Test MAE': float(np.round(mae, 4)),
    'Test R2': float(np.round(r2, 4)),
    'Precision@10': float(np.round(precision_at_k, 4)),
    'Recall@10': float(np.round(recall_at_k, 4)),
    'F1@10': float(np.round(f1_at_k, 4)),
    'NDCG@10': float(np.round(ndcg_at_k_value, 4))
}

# Convertir el diccionario a DataFrame
current_metrics_df = pd.DataFrame([current_metrics])

# Nombre del archivo CSV donde se guardarán las métricas de todos los modelos
metrics_csv = "../..//performance/model_evaluations.csv"

# Intentar cargar el CSV existente; si no existe, se crea uno nuevo
try:
    all_metrics_df = pd.read_csv(metrics_csv)
    if all_metrics_df.empty:
        all_metrics_df = pd.DataFrame(columns=current_metrics_df.columns)
except (FileNotFoundError, pd.errors.EmptyDataError):
    all_metrics_df = pd.DataFrame(columns=current_metrics_df.columns)

# Agregar la nueva entrada al DataFrame existente
all_metrics_df = pd.concat([all_metrics_df, current_metrics_df],
    ignore_index=True)

# Guardar el CSV actualizado
all_metrics_df.to_csv(metrics_csv, index=False)
print("Métricas agregadas al CSV:", metrics_csv)

```

Métricas agregadas al CSV: ../..//performance/model_evaluations.csv

/tmp/ipykernel_6095/3474256020.py:33: FutureWarning: The behavior of DataFrame concatenation with empty or all-NA entries is deprecated. In a future version, this will no longer exclude empty or all-NA columns when determining the result dtypes. To retain the old behavior, exclude the relevant entries before the concat operation.

```

    all_metrics_df = pd.concat([all_metrics_df, current_metrics_df],
    ignore_index=True)

```