

# ratings\_movies

March 28, 2025

## 0.1 Assignment 2

```
[1]: import torch
      torch.cuda.empty_cache()
```

*Import necessary libraries and load the ratings*

```
[2]: import pandas as pd
      import numpy as np

      # Cargar el archivo ratings.csv
      ratings = pd.read_csv('../data/ml-latest-small/ratings.csv')

      # Ver las primeras filas, tamaño y columnas
      print(ratings.head())
      print(ratings.shape)
      print(ratings.columns)
```

```
   userId  movieId  rating  timestamp
0        1         1     4.0  964982703
1        1         3     4.0  964981247
2        1         6     4.0  964982224
3        1        47     5.0  964983815
4        1        50     5.0  964982931
(100836, 4)
Index(['userId', 'movieId', 'rating', 'timestamp'], dtype='object')
```

*Preprocessing*

```
[3]: # Filtrar usuarios con menos de 10 ratings
      user_counts = ratings['userId'].value_counts()
      ratings = ratings[ratings['userId'].isin(user_counts[user_counts >= 10].index)]

      # Filtrar películas con menos de 10 ratings
      movie_counts = ratings['movieId'].value_counts()
      ratings = ratings[ratings['movieId'].isin(movie_counts[movie_counts >= 10].
      ↪index)]

      print(f"Usuarios después de filtrar: {ratings['userId'].nunique()}")
      print(f"Películas después de filtrar: {ratings['movieId'].nunique()}")
```

Usuarios después de filtrar: 610  
Películas después de filtrar: 2269

```
[4]: # Obtener IDs únicos
unique_user_ids = ratings['userId'].unique()
unique_movie_ids = ratings['movieId'].unique()

print(f"Número de usuarios únicos: {len(unique_user_ids)}")
print(f"Número de películas únicas: {len(unique_movie_ids)}")

# Crear diccionarios de mapeo
userId_to_index = {user_id: idx for idx, user_id in enumerate(unique_user_ids)}
movieId_to_index = {movie_id: idx for idx, movie_id in enumerate(unique_movie_ids)}

# Aplicar el mapeo al DataFrame
ratings['userIndex'] = ratings['userId'].map(userId_to_index)
ratings['movieIndex'] = ratings['movieId'].map(movieId_to_index)

# Comprobar
print(ratings.head())
```

Número de usuarios únicos: 610  
Número de películas únicas: 2269

	userId	movieId	rating	timestamp	userIndex	movieIndex
0	1	1	4.0	964982703	0	0
1	1	3	4.0	964981247	0	1
2	1	6	4.0	964982224	0	2
3	1	47	5.0	964983815	0	3
4	1	50	5.0	964982931	0	4

```
[5]: # Normalizamos ratings a [0, 1]
ratings['rating_norm'] = ratings['rating'] / 5.0
print(ratings[['rating', 'rating_norm']].head())
```

	rating	rating_norm
0	4.0	0.8
1	4.0	0.8
2	4.0	0.8
3	5.0	1.0
4	5.0	1.0

### Import Movies csv

```
[6]: # Cargar el archivo movies.csv
movies = pd.read_csv('../data/ml-latest-small/movies.csv')
print(movies.head())

# Función para obtener todos los géneros existentes
```

```

def get_all_genres(movies_df):
    genres_set = set()
    for genres in movies_df['genres']:
        for genre in genres.split("|"):
            genres_set.add(genre)
    return list(genres_set)

all_genres = get_all_genres(movies)
genre_to_index = {genre: idx for idx, genre in enumerate(all_genres)}
num_genres = len(all_genres)

# Función para codificar los géneros en un vector one-hot
def encode_genres(genres_str, genre_to_index, num_genres):
    vec = np.zeros(num_genres, dtype=np.float32)
    for genre in genres_str.split("|"):
        if genre in genre_to_index:
            vec[genre_to_index[genre]] = 1.0
    return vec

# Crear una nueva columna con el vector de géneros
movies['genres_vector'] = movies['genres'].apply(lambda x: encode_genres(x,
    ↪genre_to_index, num_genres))

```

	movieId	title \
0	1	Toy Story (1995)
1	2	Jumanji (1995)
2	3	Grumpier Old Men (1995)
3	4	Waiting to Exhale (1995)
4	5	Father of the Bride Part II (1995)

	genres
0	Adventure Animation Children Comedy Fantasy
1	Adventure Children Fantasy
2	Comedy Romance
3	Comedy Drama Romance
4	Comedy

### Combine Datasets

```

[7]: # Fusionar los datasets para agregar la información de géneros a cada rating
ratings = ratings.merge(movies[['movieId', 'genres_vector']], on='movieId',
    ↪how='left')
print(ratings.head())

```

	userId	movieId	rating	timestamp	userIndex	movieIndex	rating_norm \
0	1	1	4.0	964982703	0	0	0.8
1	1	3	4.0	964981247	0	1	0.8
2	1	6	4.0	964982224	0	2	0.8
3	1	47	5.0	964983815	0	3	1.0

4	1	50	5.0	964982931	0	4	1.0
---	---	----	-----	-----------	---	---	-----

```

                                genres_vector
0  [0.0, 0.0, 0.0, 0.0, 1.0, 0.0, 0.0, 1.0, 0.0, ...
1  [0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...
2  [0.0, 0.0, 0.0, 0.0, 0.0, 1.0, 0.0, 0.0, 0.0, ...
3  [0.0, 1.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...
4  [0.0, 1.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...

```

```

[8]: # Obtener IDs únicos y crear diccionarios de mapeo
unique_user_ids = ratings['userId'].unique()
unique_movie_ids = ratings['movieId'].unique()

print(f"Número de usuarios únicos: {len(unique_user_ids)}")
print(f"Número de películas únicas: {len(unique_movie_ids)}")

userId_to_index = {user_id: idx for idx, user_id in enumerate(unique_user_ids)}
movieId_to_index = {movie_id: idx for idx, movie_id in
    ↪ enumerate(unique_movie_ids)}

# Aplicar el mapeo al DataFrame
ratings['userIndex'] = ratings['userId'].map(userId_to_index)
ratings['movieIndex'] = ratings['movieId'].map(movieId_to_index)
print(ratings.head())

```

Número de usuarios únicos: 610

Número de películas únicas: 2269

	userId	movieId	rating	timestamp	userIndex	movieIndex	rating_norm \
0	1	1	4.0	964982703	0	0	0.8
1	1	3	4.0	964981247	0	1	0.8
2	1	6	4.0	964982224	0	2	0.8
3	1	47	5.0	964983815	0	3	1.0
4	1	50	5.0	964982931	0	4	1.0

```

                                genres_vector
0  [0.0, 0.0, 0.0, 0.0, 1.0, 0.0, 0.0, 1.0, 0.0, ...
1  [0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...
2  [0.0, 0.0, 0.0, 0.0, 0.0, 1.0, 0.0, 0.0, 0.0, ...
3  [0.0, 1.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...
4  [0.0, 1.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...

```

```

[9]: ratings['rating_norm'] = ratings['rating'] / 5.0
print(ratings[['rating', 'rating_norm']].head())

```

	rating	rating_norm
0	4.0	0.8
1	4.0	0.8
2	4.0	0.8
3	5.0	1.0

*Personalized Dataset*

```
[10]: import torch
from torch.utils.data import Dataset

class MovieLensDataset(Dataset):
    def __init__(self, users, movies, ratings, movie_features):
        self.users = users
        self.movies = movies
        self.ratings = ratings
        self.movie_features = movie_features # vector de géneros (u otras
        ↪ features)

    def __len__(self):
        return len(self.ratings)

    def __getitem__(self, idx):
        return {
            'user': self.users[idx],
            'movie': self.movies[idx],
            'rating': self.ratings[idx],
            'movie_features': self.movie_features[idx]
        }
```

*Export Dataset*

```
[11]: # Asegúrate de que 'ratings' es el DataFrame que contiene la información
        ↪ fusionada
ratings.to_csv('../data/merged_ratings_movies.csv', index=False)
print("Dataset exported successfully as 'merged_ratings_movies.csv'")
```

Dataset exported successfully as 'merged\_ratings\_movies.csv'

*Train Test Validation Split*

```
[14]: from sklearn.model_selection import train_test_split

# Filtrar usuarios con al menos 3 ratings
user_counts = ratings['userId'].value_counts()
ratings_filtered = ratings[ratings['userId'].isin(user_counts[user_counts >= 3].
        ↪ index)]

# Realizar el split por usuario
train_list = []
val_list = []
test_list = []

for user_id, group in ratings_filtered.groupby('userId'):
```

```

    user_train, user_temp = train_test_split(group, test_size=0.30,
↪random_state=42)
    user_val, user_test = train_test_split(user_temp, test_size=0.50,
↪random_state=42)
    train_list.append(user_train)
    val_list.append(user_val)
    test_list.append(user_test)

train_data = pd.concat(train_list).reset_index(drop=True)
val_data = pd.concat(val_list).reset_index(drop=True)
test_data = pd.concat(test_list).reset_index(drop=True)

print(f"Train size: {len(train_data)}")
print(f"Validation size: {len(val_data)}")
print(f"Test size: {len(test_data)}")

```

Train size: 56505

Validation size: 12164

Test size: 12447

### **Dataloaders**

```

[15]: from torch.utils.data import DataLoader
import numpy as np

# Convertir a tensores los índices y ratings (ya normalizados) como antes
train_user = torch.tensor(train_data['userIndex'].values, dtype=torch.long)
train_movie = torch.tensor(train_data['movieIndex'].values, dtype=torch.long)
train_rating = torch.tensor(train_data['rating_norm'].values, dtype=torch.
↪float32)

val_user = torch.tensor(val_data['userIndex'].values, dtype=torch.long)
val_movie = torch.tensor(val_data['movieIndex'].values, dtype=torch.long)
val_rating = torch.tensor(val_data['rating_norm'].values, dtype=torch.float32)

test_user = torch.tensor(test_data['userIndex'].values, dtype=torch.long)
test_movie = torch.tensor(test_data['movieIndex'].values, dtype=torch.long)
test_rating = torch.tensor(test_data['rating_norm'].values, dtype=torch.float32)

# Convertir la columna 'genres_vector' (obtenida al procesar movies.csv) en
↪tensores
train_movie_features = torch.tensor(np.stack(train_data['genres_vector'].
↪values), dtype=torch.float32)
val_movie_features = torch.tensor(np.stack(val_data['genres_vector'].values),
↪dtype=torch.float32)
test_movie_features = torch.tensor(np.stack(test_data['genres_vector'].values),
↪dtype=torch.float32)

```

```

# Construcción de los datasets y dataloaders
batch_size = 512

train_dataset = MovieLensDataset(train_user, train_movie, train_rating,
    ↪train_movie_features)
val_dataset = MovieLensDataset(val_user, val_movie, val_rating,
    ↪val_movie_features)
test_dataset = MovieLensDataset(test_user, test_movie, test_rating,
    ↪test_movie_features)

train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=batch_size)
test_loader = DataLoader(test_dataset, batch_size=batch_size)

```

### Model

```

[16]: import torch.nn as nn
import torch.nn.functional as F

class NeuralCollaborativeFiltering(nn.Module):
    def __init__(self, num_users, num_movies, genre_input_dim,
    ↪embedding_dim=64, genre_emb_dim=32, dropout_rate=0.3):
        super(NeuralCollaborativeFiltering, self).__init__()
        self.user_embedding = nn.Embedding(num_users, embedding_dim)
        self.movie_embedding = nn.Embedding(num_movies, embedding_dim)

        # Capa para transformar el vector de géneros a un espacio de menor
    ↪dimensión
        self.genre_layer = nn.Linear(genre_input_dim, genre_emb_dim)

        # La entrada del MLP es la concatenación de: user_embedding,
    ↪movie_embedding y genre_embedding
        input_dim = embedding_dim * 2 + genre_emb_dim
        self.fc1 = nn.Linear(input_dim, 128)
        self.fc2 = nn.Linear(128, 64)
        self.output_layer = nn.Linear(64, 1)
        self.dropout = nn.Dropout(dropout_rate)

    def forward(self, user, movie, movie_features):
        user_embedded = self.user_embedding(user)
        movie_embedded = self.movie_embedding(movie)
        genre_embedded = F.relu(self.genre_layer(movie_features))

        x = torch.cat([user_embedded, movie_embedded, genre_embedded], dim=1)
        x = F.relu(self.fc1(x))
        x = self.dropout(x)
        x = F.relu(self.fc2(x))

```

```

x = self.dropout(x)
out = self.output_layer(x)
return out.squeeze()

```

### GPU Usage

```

[17]: device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f"Usando dispositivo: {device}")

num_users = len(userId_to_index)
num_movies = len(movieId_to_index)
# 'num_genres' es el número de géneros únicos obtenido al procesar movies.csv
model = NeuralCollaborativeFiltering(num_users=num_users,
    ↪ num_movies=num_movies, genre_input_dim=num_genres).to(device)
print("Modelo con integración de features de películas creado correctamente")

```

Usando dispositivo: cuda

Modelo con integración de features de películas creado correctamente

### Loss Function and Optimizer

```

[18]: import torch.optim as optim
import torch.nn as nn

criterion = nn.MSELoss()
optimizer = optim.Adam(model.parameters(), lr=0.0005, weight_decay=1e-5)

```

### Model Training

```

[19]: num_epochs = 10 # Ajusta según convenga

for epoch in range(num_epochs):
    model.train()
    total_loss = 0
    for batch in train_loader:
        users = batch['user'].to(device)
        movies = batch['movie'].to(device)
        ratings = batch['rating'].to(device) # ya normalizado
        movie_features = batch['movie_features'].to(device)

        preds = model(users, movies, movie_features)
        loss = criterion(preds, ratings)

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        total_loss += loss.item() * len(ratings)

    avg_train_loss = total_loss / len(train_loader.dataset)

```



```

# Validación
model.eval()
val_loss = 0
with torch.no_grad():
    for batch in val_loader:
        users = batch['user'].to(device)
        movies = batch['movie'].to(device)
        ratings = batch['rating'].to(device)
        movie_features = batch['movie_features'].to(device)

        preds = model(users, movies, movie_features)
        loss = criterion(preds, ratings)
        val_loss += loss.item() * len(ratings)

    avg_val_loss = val_loss / len(val_loader.dataset)
    print(f"Epoch {epoch+1}: Train Loss = {avg_train_loss:.4f}, Val Loss = {avg_val_loss:.4f}")

```

```

Epoch 1: Train Loss = 0.1122, Val Loss = 0.0416
Epoch 2: Train Loss = 0.0588, Val Loss = 0.0389
Epoch 3: Train Loss = 0.0522, Val Loss = 0.0365
Epoch 4: Train Loss = 0.0486, Val Loss = 0.0351
Epoch 5: Train Loss = 0.0464, Val Loss = 0.0345
Epoch 6: Train Loss = 0.0450, Val Loss = 0.0341
Epoch 7: Train Loss = 0.0435, Val Loss = 0.0332
Epoch 8: Train Loss = 0.0421, Val Loss = 0.0332
Epoch 9: Train Loss = 0.0412, Val Loss = 0.0326
Epoch 10: Train Loss = 0.0404, Val Loss = 0.0328

```

### *Model Storage*

```

[ ]: torch.save(model.state_dict(), "ncf_model_current_with_movies.pth")
    print("Modelo guardado correctamente.")

```

### *Model Evaluation*

#### *RMSE*

```

[20]: all_preds = []
      all_truth = []
      model.eval()
      with torch.no_grad():
          for batch in test_loader:
              users = batch['user'].to(device)
              movies = batch['movie'].to(device)
              ratings = batch['rating'].to(device)
              movie_features = batch['movie_features'].to(device)
              preds = model(users, movies, movie_features)

```

```

        # Desnormalizamos: ratings y predicciones a escala original (por
        ejemplo, 0.5-5)
        all_preds.extend((preds * 5).cpu().numpy())
        all_truth.extend((ratings * 5).cpu().numpy())

all_preds = np.array(all_preds)
all_truth = np.array(all_truth)

rmse = np.sqrt(np.mean((all_preds - all_truth) ** 2))

```

### **MAE**

```
[21]: mae = np.mean(np.abs(all_preds - all_truth))
```

### **R-Square**

```
[22]: from sklearn.metrics import r2_score
r2 = r2_score(all_truth, all_preds)
```

### **Precision**

```
[23]: from collections import defaultdict
import numpy as np

k = 10
user_preds = defaultdict(list)
user_truth = defaultdict(list)

model.eval()
with torch.no_grad():
    for batch in test_loader:
        users = batch['user'].to(device)
        movies = batch['movie'].to(device)
        ratings = batch['rating'].to(device)
        movie_features = batch['movie_features'].to(device)
        preds = model(users, movies, movie_features)
        # Desnormalizamos: escala 0.5-5
        for u, pred, true in zip(users.cpu().numpy(), (preds * 5).cpu().
        numpy(), (ratings * 5).cpu().numpy()):
            user_preds[u].append(pred)
            user_truth[u].append(true)

precisions = []
for u in user_preds:
    preds_u = np.array(user_preds[u])
    truths_u = np.array(user_truth[u])
    # Ordenar índices según predicciones descendentes y tomar los top K
    top_k_indices = np.argsort(-preds_u)[:k]
    # Definir como relevante si el rating real es >= 4.0

```

```

    relevant = (truths_u >= 4.0)
    num_relevant = np.sum(relevant[top_k_indices])
    precision_u = num_relevant / k
    precisions.append(precision_u)

precision_at_k = np.mean(precisions)

```

### NDCG@K

```

[24]: def ndcg_at_k(relevances, k):
    relevances = np.asarray(relevances)[:k]
    if relevances.size == 0:
        return 0.0
    # DCG: usamos la fórmula  $(2^{\text{rel}} - 1) / \log_2(\text{pos} + 1)$ 
    dcg = np.sum((2 ** relevances - 1) / np.log2(np.arange(2, relevances.size + 1, 2)))
    # IDCG: DCG ideal (orden perfecto)
    ideal_relevances = np.sort(relevances)[::-1]
    idcg = np.sum((2 ** ideal_relevances - 1) / np.log2(np.arange(2, ideal_relevances.size + 2)))
    return dcg / idcg if idcg > 0 else 0.0

ndcgs = []
for u in user_preds:
    preds_u = np.array(user_preds[u])
    truths_u = np.array(user_truth[u])
    # Relevancia binaria: 1 si rating >= 4.0, 0 de lo contrario
    relevances = (truths_u >= 4.0).astype(int)
    top_k_indices = np.argsort(-preds_u)[:k]
    ndcg_u = ndcg_at_k(relevances[top_k_indices], k)
    ndcgs.append(ndcg_u)

ndcg_at_k_value = np.mean(ndcgs)

```

### All Metrics

```

[25]: from tabulate import tabulate
import pandas as pd
import numpy as np

# Suponiendo que ya tienes las variables calculadas:
# rmse, mae, r2, precision_at_k, ndcg_at_k_value

# Aseguramos que todos los valores sean Python floats y estén redondeados
metrics = {
    'RMSE': float(np.round(rmse, 4)),
    'MAE': float(np.round(mae, 4)),
    'R2': float(np.round(r2, 4)),
    'Pre': float(np.round(precision_at_k, 4)),
}

```

```

        'NDCG@10': float(np.round(ndcg_at_k_value, 4))
    }

    # Crear un DataFrame
    metrics_df = pd.DataFrame(list(metrics.items()), columns=['Métrica', 'Valor'])

    # Alternativa: formatear explícitamente los valores en el DataFrame
    metrics_df['Valor'] = metrics_df['Valor'].apply(lambda x: f"{x:.4f}")

    # Generar y mostrar la tabla con un formato más visual
    tabla_formateada = tabulate(metrics_df, headers='keys', tablefmt='pretty',
                                ↪showindex=False)
    print(tabla_formateada)

```

```

+-----+-----+
| Métrica | Valor |
+-----+-----+
| RMSE    | 0.9151 |
| MAE     | 0.7213 |
| R2      | 0.2083 |
| Pre     | 0.5048 |
| NDCG@10 | 0.8488 |
+-----+-----+

```

```

[26]: import numpy as np
import pandas as pd
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from collections import defaultdict
import torch

# Asegúrate de que model_movies está instanciado y se ha definido el
↪dispositivo 'device'
# Cargar el checkpoint (ajusta la ruta según corresponda)
model.eval()

all_preds = []
all_truth = []

with torch.no_grad():
    for batch in test_loader:
        users = batch['user'].to(device)
        movies = batch['movie'].to(device)
        ratings = batch['rating'].to(device)
        movie_features = batch['movie_features'].to(device)

        preds = model(users, movies, movie_features)

```

```

        # Desnormalizamos: suponemos que se entrenó con ratings normalizados a [0,1]
        ↪ [0,1] → multiplicamos por 5
        preds = preds * 5
        ratings = ratings * 5

        all_preds.extend(preds.cpu().numpy())
        all_truth.extend(ratings.cpu().numpy())

all_preds = np.array(all_preds)
all_truth = np.array(all_truth)

rmse = np.sqrt(mean_squared_error(all_truth, all_preds))
mae = mean_absolute_error(all_truth, all_preds)
r2 = r2_score(all_truth, all_preds)

k = 10
user_preds = defaultdict(list)
user_truth = defaultdict(list)
with torch.no_grad():
    for batch in test_loader:
        users = batch['user'].to(device)
        movies = batch['movie'].to(device)
        ratings = batch['rating'].to(device)
        movie_features = batch['movie_features'].to(device)

        preds = model(users, movies, movie_features)
        preds = preds * 5
        ratings = ratings * 5

        for u, pred, true in zip(users.cpu().numpy(), preds.cpu().numpy(),
        ↪ ratings.cpu().numpy()):
            user_preds[u].append(pred)
            user_truth[u].append(true)

precisions = []
for u in user_preds:
    preds_u = np.array(user_preds[u])
    truths_u = np.array(user_truth[u])
    top_k_indices = np.argsort(-preds_u)[:k]
    relevant = (truths_u >= 4.0)
    num_relevant = np.sum(relevant[top_k_indices])
    precisions.append(num_relevant / k)
precision_at_k = np.mean(precisions)

def ndcg_at_k(relevances, k):
    relevances = np.asarray(relevances)[:k]
    if relevances.size == 0:

```

```

        return 0.0
    dcg = np.sum((2**relevances - 1) / np.log2(np.arange(2, relevances.size +
↪2)))
    ideal_relevances = np.sort(relevances)[::-1]
    idcg = np.sum((2**ideal_relevances - 1) / np.log2(np.arange(2,
↪ideal_relevances.size + 2)))
    return dcg / idcg if idcg > 0 else 0.0

ndcgs = []
for u in user_preds:
    preds_u = np.array(user_preds[u])
    truths_u = np.array(user_truth[u])
    relevances = (truths_u >= 4.0).astype(int)
    top_k_indices = np.argsort(-preds_u)[:k]
    ndcg_u = ndcg_at_k(relevances[top_k_indices], k)
    ndcgs.append(ndcg_u)
ndcg_at_k_value = np.mean(ndcgs)

# Crear el diccionario de métricas y exportar a CSV
metrics = {
    "Model": "NCFMovies (Ratings + Genres)",
    "Test RMSE": rmse,
    "Test MAE": mae,
    "Test R2": r2,
    "Precision@10": precision_at_k,
    "NDCG@10": ndcg_at_k_value
}
metrics_df = pd.DataFrame([metrics])
metrics_df.to_csv("ncf_movies_metrics.csv", index=False)
print("Model 2 metrics exported to 'ncf_movies_metrics.csv'")

```

Model 2 metrics exported to 'ncf\_movies\_metrics.csv'